# Memory Allocation: Either Love It or Hate It (or just think it's okay)

Andrei Alexandrescu

© 2007-2008 Andrei Alexandrescu 1 / 45 1

# **Memory Allocation**

- Fundamental part of most systems since 1960
- Frequently at the top of time histograms
- Hard to improve on
- Two views:
  - Solved
  - Insoluble

© 2007-2008 Andrei Alexandrescu 2 / 45

### **Overview**

Lies, Damn Lies...

**Custom Allocators** 

**Policy-based Implementation** 

**Conclusions** 

© 2007-2008 Andrei Alexandrescu 3 / 45

Memory Allocation

#### Lies, Damn Lies...

- Memory Allocation
- Engineering

#### Challenges

Capacity and

#### Generalization

- Odysseus, Beware of Sirens
- The Windows XP

#### Allocator

- Windows XP-Specific Heap
- The Doug Lea Allocator
- Benchmarks
- Results

#### **Custom Allocators**

Policy-based Implementation

Conclusions

Lies, Damn Lies...

© 2007-2008 Andrei Alexandrescu 4 / 45

# **Memory Allocation**

Need: randomly allocate and free chunks of various sizes

Fact: most applications have highly regular allocation/deallocation patterns

 Challenge: Such patterns notoriously escape standard statistical analysis

© 2007-2008 Andrei Alexandrescu 5 / 4

# **Engineering Challenges**

- Highly competitive:
  - General-purpose allocators: 100 cycles/alloc
  - Specialized allocators: <12 cycles/alloc</li>
- Hard to balance capacity with generalization
  - "Too general:" mediocre for all patterns
  - "Overspecialized:" numerous degenerate situations
- Very hard to modularize/componentize

c) 2007-2008 Andrei Alexandrescu

# **Capacity and Generalization**

- Generalization: Figure out patterns from seen data
  - Too much: "Rectangular piece of paper? That's a \$100 bill"
- Capacity: Ability to allocate new categories for data
  - Too much: "#L26118670? It's a fake; all \$100 bills
     I've seen had other serial numbers"
- They are competitive with one another
- How to strike the right balance?

 $\odot$  2007-2008 Andrei Alexandrescu 7 / 4

# **Odysseus, Beware of Sirens**

- Myth: application efficiency can be largely increased by hand-writing a custom memory allocator
- Fact: on six out of eight popular benchmarks, replacing back the custom allocator with a general-purpose allocator generally improves speed and memory consumption [1]
- Fact: simplest custom allocators tuned for special situations are the only ones providing a measurable gain

© 2007-2008 Andrei Alexandrescu 8 / 4

### The Windows XP Allocator

- Best-fit allocator
- 127 free lists for 8, 16, 24, 32..., 1024 bytes
- >1024 bytes from sorted linked list
  - Sacrifices speed for size accuracy

© 2007-2008 Andrei Alexandrescu 9 / 45

# **Windows XP-Specific Heap**

- Very rich interface
  - Multiple heaps
  - Region semantics
  - Individual Object Deletion
- Distinctly unimpressive performance

© 2007-2008 Andrei Alexandrescu 10 / 45

# The Doug Lea Allocator

- Approximate best-fit allocator
- Quicklists for <64-bytes allocations</li>
- Coalesce and split blocks on medium-sized allocations (<128KB)</li>
- Large blocks allocated with mmap

- The basis of the Linux allocator
- Honed throughout millennia

© 2007-2008 Andrei Alexandrescu 11 / 43

# **Benchmarks**

#	Name	Description
1.	197.parser	English parser
2.	boxed-sim	Balls-in-box simulator
3.	C-breeze	C-to-C optimizing compiler
4.	175.vpr	FPGA placement and routing
5.	176.gcc	Optimizing C compiler
6.	apache	Web server
7.	lcc	Retargetable C compiler
8.	muddle	MUD compiler/interpreter

© 2007-2008 Andrei Alexandrescu 12 / 45

### Results

- Pretty much everybody beats the Windows XP allocator in speed
  - Space consumed not measurable

 Only 1cc and muddle outperformed Doug Lea's allocator (20-47% in speed, 25% in space)

Rest: within 10% in speed, but consuming a hefty 33% extra space

© 2007-2008 Andrei Alexandrescu 13 / 4

Memory Allocation

Lies, Damn Lies...

#### **Custom Allocators**

Temporal Allocation

#### Patterns

- Size Distribution
- Useful Custom

#### Allocators

- Regions
- How Regions Work
- Free Lists
- How Free Lists Work

Policy-based Implementation

Conclusions

### **Custom Allocators**

© 2007-2008 Andrei Alexandrescu 14 / 45

# **Temporal Allocation Patterns**

- Ramps: Slowly chews memory (logs, editors, file processing, debuggers, leaks)
- Plateaus: Dynamic equilibrium—the application statistically allocates as much as it deallocates
- Peaks: The application suddenly allocates much memory and then deallocates most of it (compilers, linkers, parsers, other staged tools)

© 2007-2008 Andrei Alexandrescu 15 / 45

### **Size Distribution**

- Pareto-distributed (80/20): Most allocations are of a small set of sizes
  - Focus on allocating specific sizes

- Long-tail: Distribution does not have a strong bias
  - Focus on allocating random sizes

© 2007-2008 Andrei Alexandrescu 16 / 45

### **Useful Custom Allocators**

- Mostly there are two effective custom allocators
  - Regions (a.k.a. arenas, pools)
  - Free Lists
- Both only work for particular allocation patterns
- Both may do very, very bad on other allocation patterns
- Precisely because everything that works for general allocation patterns has been already incorporated in today's allocators!

© 2007-2008 Andrei Alexandrescu 17 / 49

# Regions

- Allocate memory blocks of arbitrary sizes
- Deallocate all at once
- Useful for batch, transaction-oriented applications
  - Web servers
  - Compilers
  - Databases
- API:

```
void * allocate(size_t size);
void freeAll();
```

© 2007-2008 Andrei Alexandrescu 18 / 45

# **How Regions Work**

- Initially allocate a large block of memory
- Each allocation is a bump of a pointer
  - Plus checking for overrun
- There's no deallocation
- Terminating the region deallocates the large block

 Risk: pointers into the region from global heap or longer-lived regions

© 2007-2008 Andrei Alexandrescu 19 / 45

### **Free Lists**

- Useful for Pareto-distributed block sizes
  - Stock data
  - String processing apps
  - Scientific apps
- API:

```
void * allocate(size_t size);
void deallocate(void * p);
size_t allocatedSize(void * p);
```

© 2007-2008 Andrei Alexandrescu 20 / 45

### **How Free Lists Work**

- Allocate memory blocks of arbitrary sizes
- Remember blocks of specific size in a list
- Subsequent allocations of that size are O(1)
- The free list can live inside the freed blocks

- Advantage: superfast for one special size
- Risk: fixed size, no coallescing, no reallocation, extra space consumed

© 2007-2008 Andrei Alexandrescu 21 / 45

Memory Allocation

Lies, Damn Lies...

**Custom Allocators** 

#### Policy-based Implementation

- General Interface
- Top Allocator
- ... continued
- Supporting the

Unsupported

- ... continued
- In-Situ Regions
- In-Situ Regions
- Heap Regions
- Lists of Allocators
- Policy-Based Design:

Yum (I)

- Free Lists
- Free Lists (cont'd)
- Aside
- Aside (cont'd)
- Rounded Free Lists
- Example:

Ghostscript [3]

- Policy-Based Design: Yum (II)
- tuiii (ii
- ReapsPolicy-Based Design:

Yum (III)

# Policy-based Implementation

### **General Interface**

```
template <class BaseAllocator>
struct Allocator : private BaseAllocator {
    Allocator();
    ~Allocator();
    void * allocate(size_t size);
    void deallocate(void * p);
    size_t allocatedSize(void * p);
}
```

- "For adoption" structure ⇒ efficient layering
- allocatedSize is optional
- allocatedSize \geq requested

© 2007-2008 Andrei Alexandrescu 23 / 4

# **Top Allocator**

```
struct Mallocator {
    static void * allocate(size_t size) {
        return malloc(size);
    }
    void deallocate(void * p) {
        free(p);
    }
```

© 2007-2008 Andrei Alexandrescu 24 / 45

#### ... continued

```
#if defined(WIN32)
   size_t allocatedSize(void * p) {
      return _msize(p);
#elif defined(GNU)
   size_t allocatedSize(void * p) {
      return malloc_usable_size(p);
#endif
};
```

© 2007-2008 Andrei Alexandrescu 25 / 4

# **Supporting the Unsupported**

```
template <class B>
struct SizedAlloc : private B {
   void * allocate(size_t s) {
      size_t * pS = static_cast<size_t*>(
         B::allocate(s + sizeof(size_t)));
      return *pS = s, pS + 1;
   void deallocate(void* p) {
      B::deallocate(
         static_cast<size_t*>(p) - 1);
```

© 2007-2008 Andrei Alexandrescu 26 / 4

#### ... continued

```
size_t allocatedSize(void * p) {
    return (static_cast<size_t*>(p))[-1];
}
};
```

- Remark: Add portable size computation over any allocator, not only malloc
  - We'll make use of that in a few slides

© 2007-2008 Andrei Alexandrescu 27 / 4

# **In-Situ Regions**

```
template <size_t S> struct InSituRegion {
   void * allocate(size_t s) {
      s = (s + M) & ^{\sim}M;
      if (s > S \mid | p > buf + S - s) return 0;
      void * r = p;
      p += s;
      return r;
   enum { A=2, M=(1<<A)-1}; // align
   private: unsigned char buf[S], *p;
};
```

© 2007-2008 Andrei Alexandrescu 28 / 4

# **In-Situ Regions**

```
template <size_t S> struct InSituRegion {
    ...
    void freeAll() {
        p = buf;
    }
    ...
};
```

- No allocatedSize
- Going out of scope entails an unchecked freeAll

© 2007-2008 Andrei Alexandrescu 29 / 4

# **Heap Regions**

```
template <class B>
struct HeapRegion : private B {
   HeapRegion(size_t s = 1024 * 128)  {
      buf = p = enforce(B::allocate(s));
      limit = buf + s;
   ~HeapRegion() {
      B::deallocate(buf);
   private: unsigned char *buf, *limit, *p;
};
```

© 2007-2008 Andrei Alexandrescu 30 / 4

### **Lists of Allocators**

- How to make regions no-fail?
- Keep a list of regions
- Allocate from head of the list
- Add a new element when head full

```
template <class A> struct AllocList {
    ...
    private: slist<A> allocs;
};
```

Works both with in-situ and heap allocators

© 2007-2008 Andrei Alexandrescu 31 / 45

# Policy-Based Design: Yum (I)

typedef AllocList<InSituRegion<1024 \* 128> >
 ScalableRegion;

© 2007-2008 Andrei Alexandrescu 32 / 4

#### **Free Lists**

```
template <size_t S, class B>
struct ExactFreeList : private B {
   void * allocate(size_t s) {
      if (s != S || !list)
         return B::allocate(s);
      void * r = list;
      list = list->next;
      return r;
```

© 2007-2008 Andrei Alexandrescu 33 / 4

# Free Lists (cont'd)

```
void deallocate(void * p) {
      if (allocatedSize(p) != S)
         return B::deallocate(p);
      list * pL = static_cast<List*>(p);
      pL->next_ = list_;
      list_= pL;
   using B::allocatedSize;
private:
   struct List { List * next; } list;
};
```

© 2007-2008 Andrei Alexandrescu 34 / 4

### **Aside**

• using: the enemy of generic programmers

```
struct A {
    ... no declaration of allocatedSize ...
};
template <class B>
struct C : private B {
    using B::allocatedSize;
}
C<A> object; // error!
```

© 2007-2008 Andrei Alexandrescu 35 / 4

# Aside (cont'd)

However, this does work:

```
struct A {
   ... no declaration of allocatedSize ...
};
template <class B>
struct C : private B {
   size_t allocatedSize() {
      return B::allocatedSize();
C<A> object; // fine
```

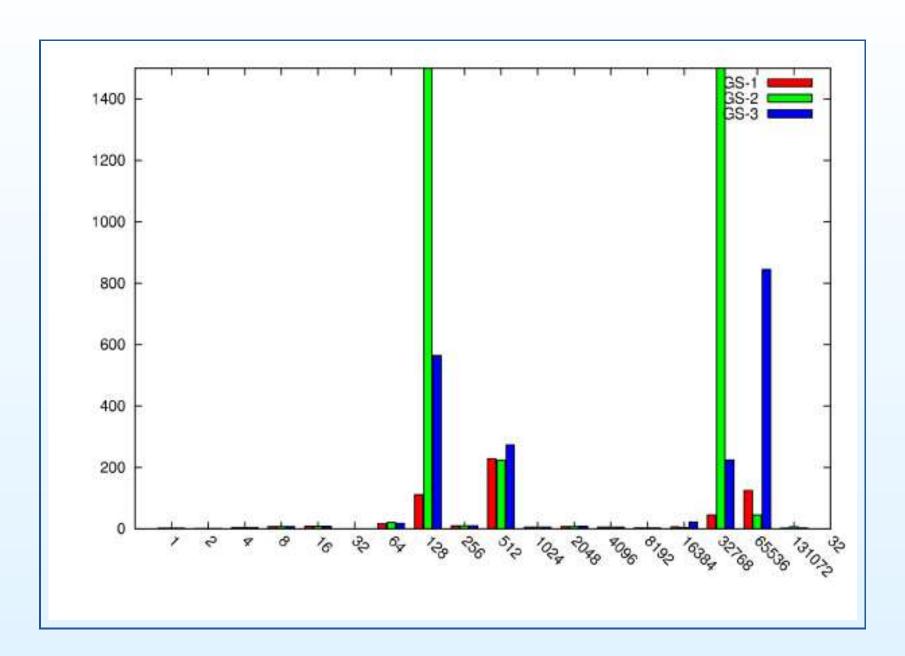
© 2007-2008 Andrei Alexandrescu 36 / 49

#### **Rounded Free Lists**

```
template <class B>
struct RoundedFreeList : private B {
   RoundedFreeList(size_t preferred) {
      list = B::allocate(preferred);
      list->next = 0;
      S = allocSize(list);
   private: size_t S;
};
```

© 2007-2008 Andrei Alexandrescu 37 / 49

# **Example: Ghostscript [3]**



© 2007-2008 Andrei Alexandrescu 38 / 45

# Policy-Based Design: Yum (II)

© 2007-2008 Andrei Alexandrescu 39 / 4

# Reaps

- Regions can't deallocate (deallocate is a no-op)
- Free lists intercept deallocate calls to remember deallocated blocks of a specific size

• ... Wait a **minute**!

© 2007-2008 Andrei Alexandrescu 40 / 45

# Policy-Based Design: Yum (III)

- Allocate from the region (fast)
- Part of what's deallocated gets under free list control (fast)
- Combine benefits of regions and free lists!

© 2007-2008 Andrei Alexandrescu 41 / 4

# **General Reaps**

- Berger's insight [1]:
  - Any pointer freed from a reap can be passed to another memory manager
  - Best-fit, free lists, quick lists...
  - Complication: some of these might need to add metadata to each allocation
- Modular implementation in HeapLayers [2]
- Yet competitive with hand-written special allocators

© 2007-2008 Andrei Alexandrescu 42 / 45

Memory Allocation

Lies, Damn Lies...

**Custom Allocators** 

Policy-based Implementation

#### Conclusions

- To Conclude
- References

# **Conclusions**

© 2007-2008 Andrei Alexandrescu 43 / 45

#### **To Conclude**

To speed up applications:

```
if (profilerClearlyShowsAllocBottleneck()) {
   if (needToUseRegions()) {
     useReaps();
   } else {
     assert(!useWindowsAllocator());
     useDlMalloc();
   }
}
```

© 2007-2008 Andrei Alexandrescu 44 / 4

#### References

Memory Allocation

Lies, Damn Lies...

**Custom Allocators** 

Policy-based Implementation

#### Conclusions

- To Conclude
- References

- [1] E. Belkin, B. Zorn, and K. McKinley. Reconsidering Custom Memory Allocation. OOPSLA, 2002.
- [2] Emery D. Berger et al. Composing High-Performance Memory Allocators. In SIGPLAN, pages 114–124, 2001.
- [3] Miguel Masmano. Histograms Library. http://tinyurl.com/2n5nds.

© 2007-2008 Andrei Alexandrescu 45 / 45